

UDARI MADHUSHANI

Ph.D. Candidate
Princeton University, Princeton, NJ 08544

☎ (609)-933-7329
🌐 [udarimadhu.github.io](https://github.com/udarimadhu) ✉ udarim@princeton.edu

RESEARCH INTEREST

Human-AI coordination, Multi-agent reinforcement learning (communication, zero-shot generalization, evolution of cooperation), Game Theory, Deep reinforcement learning, Sequential decision making, Bandit Theory, Robotics, Control theory.

EDUCATION

Program	Institution	Years
Ph.D., Mechanical Engineering	Princeton University NJ, USA	2017 - Present
M.A., Mechanical Engineering	Princeton University NJ, USA	2017 - 2019
B.Sc. Eng (Hons), Electronic Engineering	University of Peradeniya Kandy, Sri Lanka	2011 - 2015

HONORS

Graduate Honors (Princeton University)

- Harold W. Dodds Fellowship (prestigious honorific fellowship from Graduate School) 2021-22
- Britt and Eli Harari Fellowship 2020-21
- Larisse Rosentweig Klein Memorial Award Fall 2019
- Martin Summerfield Graduate Fellowship 2018-19
- Athena-Feron Prize for Mathematical Excellence Fall 2018
- Elliotte Robinson Little '25 Student Aid Fund Fellowship 2017-18

Other Awards

- Presidential Award for Scientific Research (Sri Lanka) 2019
- Best Undergraduate Project Award (IEEE Sri Lanka Section) December 2015

EXPERIENCE

Research Internships

- Google Deepmind, London, UK May 2022 - Sep 2022
Mentors: Daniel Hennes and Edgar Duenez-Guzman
Project: Zero-shot Generalization in Mixed Motive Games
- Facebook AI Research, Menlo Park, California, USA May 2021 - August 2021
Mentors: Kalesha Bullard and Roberto Calandra
Project: Incentivizing Coordination in Multi-Agent Reinforcement Learning
- Siemens, Princeton, New Jersey, USA May 2020 - August 2020
Mentors: Biswadip Dey and Amit Chakraborty
Project: Using Hamiltonian Monte Carlo Sampling for Reinforcement Learning Problems in High-dimension

Assistant in Research

- Princeton University September 2017 - Present
- Sri Lanka Technological Campus (SLTC), Sri Lanka January 2016 - August 2017

Assistant in Teaching

- MAE 542 Advanced Dynamics (Princeton University) Fall 2020

- MAE 502/APC 506 Mathematical Methods of Engineering Analysis (Princeton University) Spring 2020
- MAE 345/MAE 549 Introduction to Robotics (Princeton University) Fall 2019
- EE 554 Microwave Techniques (University of Peradeniya), Sri Lanka October 2015 - December 2015

OUTREACH

Mentoring

Princeton University

- Justin Lidard (Ph.D. Mechanical Engineering) 2021-Present
- Kathryn Wantlin (M.Sc. Computer Science) 2021-Present
- Sarah Dillender (B.S.E. Mechanical Engineering) 2020-21
Thesis: Message passing structures for improved policy finding in decentralized multi-agent Q-learning.
- Gargi Sadalgekar, Samarie Wilson and Jacob Walrath (B.S.E. Mechanical Engineering) 2020-21
Thesis: Decision making and task allocation in a multi-robot system.

Peradeniya University, Sri Lanka

- Lasitha Weerakon (Post graduate, Mechanical Engineering) 2016-17
Published in American Control Conference (ACC), 2018.
- Isuru Basnayake (Post graduate, Mechanical Engineering) 2016-17
Published in International Conference of Industrial and Information Systems, (ICIIS), 2017.
- Kusal Tennakoon (Post graduate, Mechanical Engineering) 2016-17
Published in International Conference of Industrial and Information Systems, (ICIIS), 2017.

Commitment to Diversity and Inclusion

- Co-organized "Re-education Monthly Book Club" aimed towards improving awareness of implicit biases (Princeton University) 2020-Present
- Co-organized the workshop "Diversity and Mentorship" at American Control Conference May 2021
- Co-organized the workshop "Inclusive Teaching" (Princeton University) Jan 2021

Leadership Roles

- Committee member of Women in STEM Leadership Council (Princeton University) 2020-Present
- Committee member of Graduate Women in Science & Engineering (Princeton University) 2020-Present
- Future Digileaders Training (Stockholm, Sweden) November 2019
- Founder of EE Instructors Scholarship (Peradeniya University, Sri Lanka) 2016

Volunteer for

- World Maker Faire September 2018
- American Control Conference June 2018

ACADEMIC SERVICES

Program committee Member of

- Adaptive and Learning Agents Workshop (**ALA**) at International Conference on Autonomous Agents and Multi Agent Systems (**AAMAS**) 2022
- Cooperative AI Workshop at Conference on Neural Information Processing Systems (**NeurIPS**) 2021

Co-organizer of

- Cooperative AI Workshop at International Conference on Machine Learning (**ICML**) 2023 (proposal in preparation)
- Special issue on Evolutionary Dynamics of Collective Human-Artificial Intelligence at Proceedings of National Academy of Science (**PNAS**) 2023

Reviewer for

Journal: Automatica, IEEE Transactions on Automatic Control, IEEE Control Systems Letters, Journal of Field Robotics

Conference: Conference on Neural Information Processing Systems, International Conference on Machine Learning, IEEE Conference on Decision and Control, American Control Conference, International Conference on Intelligent Robots and Systems, European Control Conference, International Conference of Industrial and Information Systems

Student Member of

IEEE (Institute of Electrical and Electronics Engineers), IEEE Young Professionals, IEEE Women in Engineering, APS (American Physical Society)

PUBLICATIONS

In Preparation

- [P3] Zero-shot Generalization in Mixed Motive Games
Udari Madhushani, John Agapiou, Kevin McKee, Joel Z Leibo, Thomas Anthony, Richard Everett and Edgar Duéñez-Guzmán
(Work done while interning at Deepmind)
- [P2] Unilateral Incentive Alignment in Two-player Stochastic Games
Alex McAvoy, **Udari Madhushani**, Christian Hilbe, Krishnendu Chatterjee, Joshua B. Plotkin
(Proceedings of National Academy of Science - special issue on Collective Artificial Intelligence and Evolutionary Dynamics)
- [P1] Collective Cooperative Intelligence
Wolfram Barfuss, Lewis Hammond, Christian Hilbe, Joel Leibo, Chaitanya Gokhale, **Udari Madhushani**, Edward Hughes, Tom Lenaerts
(Proceedings of National Academy of Science - special issue on Collective Artificial Intelligence and Evolutionary Dynamics)

Technical Reports

- [T1] **Melting Pot 2.0**
John Agapiou, Alexander Vezhnevets, Duéñez-Guzmán, Jayd Matyas, Yiran Mao, Peter Sunehag, Raphael Köster, **Udari Madhushani**, Kavya Kopparapu, Ramona Comanescu, DJ Strouse, Michael B. Johanson, Sukhdeep Singh, Julia Haas, Igor Mordatch, Dean Mobbs, Joel Z Leibo

Journal Papers

- [J2] **Heterogeneous Explore-Exploit Strategies on Multi-Star Networks**
Udari Madhushani, Naomi Leonard
IEEE Control Systems Letters, 2021.
- [J1] **Semi-globally Exponential Trajectory Tracking for a Class of Spherical Robots**
Udari Madhushani, Sanjeeva Maithripala, Janaka Wijayakulasooriya, Jordan Berg
Automatica, 2017.

Peer-reviewed Conference Papers

- [C13] On Using Hamiltonian Monte Carlo Sampling for Reinforcement Learning
Udari Madhushani, Biswadip Dey, Naomi Leonard, Amit Chakraborty
Conference on Decision and Control, **(CDC)** 2022. **(Invited Paper)**
- [C12] **A Regret Minimization Approach to Multi-Agent Control**
Udaya Ghai, **Udari Madhushani**, Naomi Leonard, Elad Hazan
International Conference on Machine Learning, **(ICML)** 2022.
- [C11] **Provably Efficient Multi-Agent Reinforcement Learning with Fully Decentralized Communication**
Justin Lidard, **Udari Madhushani**, Naomi Leonard
American Control Conference, **(ACC)** 2022.
- [C10] **One More Step Towards Reality: Cooperative Bandits with Imperfect Communication**
Udari Madhushani, Abhimanyu Dubey, Naomi Leonard, Alex Pentland
Conference on Neural Information Processing Systems, **(NeurIPS)** 2021.
- [C9] **Multi-robot Learning and Coverage of Unknown Fields**
Maria Santos, **Udari Madhushani**, Alessia Benevento, Naomi Leonard
IEEE International Symposium on Multi-Robot and Multi-Agent Systems, **(MRS)** 2021.

- [C8] **Distributed Bandits: Probabilistic Communication on d -regular Graphs**
Udari Madhushani, Naomi Leonard
 European Control Conference, (**ECC**) 2021.
- [C7] **A Dynamic Observation Strategy for Multi-agent Multi-armed Bandit Problem**
Udari Madhushani, Naomi Leonard
 European Control Conference, (**ECC**) 2020.
- [C6] **Heterogeneous Stochastic Interactions for Multiple Agents in a Multi-armed Bandit Problem**
Udari Madhushani, Naomi Leonard
 European Control Conference, (**ECC**) 2019. (**Invited Paper**)
- [C5] **Feedback Regularization and Geometric PID Control for Robust Stabilization of a Planar Three-link Hybrid Bipedal Walking Model**
 Lasitha Weerakon, **Udari Madhushani**, Sanjeeva Mathripala, Jordan Berg
 American Control Conference, (**ACC**) 2018.
- [C4] **Multi-armed Bandit Based Approach for Performance Enhancement of Window Intensity Test (WIT) Detector**
 Kusal Tennakoon, **Udari Madhushani**, Sanjeeva Mathripala
 International Conference of Industrial and Information Systems, (**ICIIS**) 2017.
- [C3] **Intrinsic PID Controller for a Segway Type Mobile Robot**
 Isuru Basnayake, **Udari Madhushani**, Sanjeeva Mathripala
 International Conference of Industrial and Information Systems, (**ICIIS**) 2017.
- [C2] **Feedback Regularization and Geometric PID Control for Trajectory Tracking of Mechanical Systems: Hoop Robots on an Inclined Plane**
Udari Madhushani, Sanjeeva Mathripala, Jordan Berg
 American Control Conference, (**ACC**) 2017.
- [C1] **WIT: Window Intensity Test Detector and Descriptor**
Udari Madhushani, Sanjeeva Mathripala, Janaka Wijayakulasooriya
 International Conference of Industrial and Information Systems, (**ICIIS**) 2016.

Peer-reviewed Workshop Papers

- [W7] **Multi-robot Learning and Coverage of Unknown Fields**
 Maria Santos, **Udari Madhushani**, Alessia Benevento, Naomi Leonard
 Autonomous Robots and Multirobot Systems (**ARMS**), International Conference on Autonomous Agents and Multiagent Systems (**AAMAS**) 2022.
- [W6] **A Regret Minimization Approach to Multi-Agent Control**
 Udaya Ghai, **Udari Madhushani**, Naomi Leonard, Elad Hazan
 Workshop on Gamification and Multiagent Solutions (**GMS**), International Conference on Learning Representations, (**ICLR**) 2022. (**Oral - Won the best poster award**)
- [W5] **Provably Efficient Decentralized Communication for Multi-Agent RL**
 Justin Lidard, **Udari Madhushani**, Naomi Leonard
 Reinforcement Learning in Games (**RLG**), AAAI Conference on Artificial Intelligence, (**AAAI**) 2022. (**Oral**)
- [W4] **On Using Hamiltonian Monte Carlo Sampling for Reinforcement Learning Problems in High-dimension**
Udari Madhushani, Biswadip Dey, Naomi Leonard, Amit Chakraborty
 Deep Reinforcement Learning Workshop (**Deep RL**), Conference on Neural Information Processing Systems, (**NeurIPS**) 2021.
- [W3] **When to Call Your Neighbor? Strategic Communication in Cooperative Stochastic Bandits**
Udari Madhushani, Naomi Leonard
 Learning and Decision-Making with Strategic Feedback (**StratML**), Conference on Neural Information Processing Systems, (**NeurIPS**) 2021.

- [W2] **It Doesn't Get Better and Here's Why: A Fundamental Drawback in Natural Extensions of UCB to Multi-agent Bandits**
Udari Madhushani, Naomi Leonard
I can't Believe It's Not Better! Workshop (**ICBINB**), Conference on Neural Information Processing Systems, (**NeurIPS**) 2020. (**Spotlight**)
- [W1] **Distributed Learning: Sequential Decision Making in Resource-Constrained Environments**
Udari Madhushani, Naomi Leonard
Practical ML for Developing Countries Workshop (**PML4DC**), International Conference on Learning Representations, (**ICLR**) 2020. (**Oral**)

SKILLS

Computer Skills: Python, PyTorch, TensorFlow, JAX, Matlab, \LaTeX

Language Skills: Sinhala (Native), English (Fluent)